

# Machine Learning Semester Project

NAME: SAYAB ABBASI

ID: 37182

## Article NAME: Patient Diet Recommendation System through Machine Learning Model

### DATASET NAME: Personalized Medical Diet Recommendations Dataset

In [600]:

```
import pandas as pd
import numpy as np
from collections import Counter
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler, LabelEncoder
from sklearn.metrics import accuracy_score, classification_report
from sklearn.multioutput import MultiOutputClassifier
from xgboost import XGBClassifier
from imblearn.over_sampling import SMOTE
```

## EDA

In [601]:

```
df = pd.read_csv("/content/Personalized_Diet_Recommendations.csv")
```

In [602]:

```
df
```

Out[602]:

	Patient_ID	Age	Gender	Height_cm	Weight_kg	BMI	Chronic_Disease	Blood_Pressure_Systolic	Blood_Pressure_Diast
0	P00001	56	Other	163	66	24.84	NaN	175	
1	P00002	69	Female	171	114	38.99	NaN	155	
2	P00003	46	Female	172	119	40.22	NaN	137	
3	P00004	32	Female	197	118	30.41	NaN	148	
4	P00005	60	Female	156	109	44.79	Hypertension	160	
...	...	...	...	...	...	...	...	...	
4995	P04996	42	Female	172	99	33.46	NaN	115	
4996	P04997	39	Female	155	61	25.39	NaN	110	
4997	P04998	48	Female	165	61	22.41	Diabetes	113	
4998	P04999	34	Other	151	82	35.96	Heart Disease	105	
4999	P05000	72	Other	173	98	32.74	NaN	121	

5000 rows x 30 columns

In [603]:

```
df.head()
```

Out[603]:

	Patient_ID	Age	Gender	Height_cm	Weight_kg	BMI	Chronic_Disease	Blood_Pressure_Systolic	Blood_Pressure_Diastolic
0	P00001	56	Other	163	66	24.84	NaN	175	75
1	P00002	69	Female	171	114	38.99	NaN	155	72
2	P00003	46	Female	172	119	40.22	NaN	137	101
3	P00004	32	Female	197	118	30.41	NaN	148	91
4	P00005	60	Female	156	109	44.79	Hypertension	160	109

5 rows x 30 columns



In [604]:

```
df.tail()
```

Out[604]:

	Patient_ID	Age	Gender	Height_cm	Weight_kg	BMI	Chronic_Disease	Blood_Pressure_Systolic	Blood_Pressure_Diast
4995	P04996	42	Female	172	99	33.46	NaN	115	
4996	P04997	39	Female	155	61	25.39	NaN	110	
4997	P04998	48	Female	165	61	22.41	Diabetes	113	
4998	P04999	34	Other	151	82	35.96	Heart Disease	105	
4999	P05000	72	Other	173	98	32.74	NaN	121	

5 rows x 30 columns



In [605]:

```
df.describe()
```

Out[605]:

	Age	Height_cm	Weight_kg	BMI	Blood_Pressure_Systolic	Blood_Pressure_Diastolic	Cholesterol_Lev
count	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000	5000.000000
mean	48.805600	174.244000	84.36620	28.353134	133.982400	89.735800	224.297800
std	17.906991	14.229173	20.18103	8.297745	26.216215	17.283025	42.918900
min	18.000000	150.000000	50.00000	12.630000	90.000000	60.000000	150.000000
25%	34.000000	162.000000	67.00000	21.850000	111.000000	75.000000	187.000000
50%	49.000000	174.000000	84.00000	27.640000	133.000000	90.000000	224.000000
75%	64.000000	186.000000	102.00000	33.812500	157.000000	105.000000	261.000000
max	79.000000	199.000000	119.00000	52.890000	179.000000	119.000000	299.000000



In [606]:

```
df.info()
```

<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 5000 entries, 0 to 4999  
Data columns (total 30 columns):  
# Column Non-Null Count Dtype

```
-----
0 Patient_ID 5000 non-null object
1 Age 5000 non-null int64
2 Gender 5000 non-null object
3 Height_cm 5000 non-null int64
4 Weight_kg 5000 non-null int64
5 BMI 5000 non-null float64
6 Chronic_Disease 2957 non-null object
7 Blood_Pressure_Systolic 5000 non-null int64
8 Blood_Pressure_Diastolic 5000 non-null int64
9 Cholesterol_Level 5000 non-null int64
10 Blood_Sugar_Level 5000 non-null int64
11 Genetic_Risk_Factor 5000 non-null object
12 Allergies 1503 non-null object
13 Daily_Steps 5000 non-null int64
14 Exercise_Frequency 5000 non-null int64
15 Sleep_Hours 5000 non-null float64
16 Alcohol_Consumption 5000 non-null object
17 Smoking_Habit 5000 non-null object
18 Dietary_Habits 5000 non-null object
19 Caloric_Intake 5000 non-null int64
20 Protein_Intake 5000 non-null int64
21 Carbohydrate_Intake 5000 non-null int64
22 Fat_Intake 5000 non-null int64
23 Preferred_Cuisine 5000 non-null object
24 Food_Aversions 3775 non-null object
25 Recommended_Calories 5000 non-null int64
26 Recommended_Protein 5000 non-null int64
27 Recommended_Carbs 5000 non-null int64
28 Recommended_Fats 5000 non-null int64
29 Recommended_Meal_Plan 5000 non-null object
dtypes: float64(2), int64(17), object(11)
memory usage: 1.1+ MB
```

In [607]:

```
df.isnull().sum()
```

Out[607]:

	0
<b>Patient_ID</b>	0
<b>Age</b>	0
<b>Gender</b>	0
<b>Height_cm</b>	0
<b>Weight_kg</b>	0
<b>BMI</b>	0
<b>Chronic_Disease</b>	2043
<b>Blood_Pressure_Systolic</b>	0
<b>Blood_Pressure_Diastolic</b>	0
<b>Cholesterol_Level</b>	0
<b>Blood_Sugar_Level</b>	0
<b>Genetic_Risk_Factor</b>	0
<b>Allergies</b>	3497
<b>Daily_Steps</b>	0
<b>Exercise_Frequency</b>	0
<b>Sleep_Hours</b>	0
<b>Alcohol_Consumption</b>	0
<b>Smoking_Habit</b>	0
<b>Dietary_Habits</b>	0

Caloric_Intake	0
Protein_Intake	0
Carbohydrate_Intake	0
Fat_Intake	0
Preferred_Cuisine	0
Food_Aversions	1225
Recommended_Calories	0
Recommended_Protein	0
Recommended_Carbs	0
Recommended_Fats	0
Recommended_Meal_Plan	0

dtype: int64

## Data Set Summry

Type	Examples	Description
Demographics	Age , Gender , Height_cm , Weight_kg	Basic patient info
Health Indicators	BMI , Blood_Pressure , Cholesterol_Level , Blood_Sugar_Level	Key medical metrics
Medical History	Chronic_Disease , Genetic_Risk_Factor , Allergies	Existing conditions and risks
Lifestyle	Daily_Steps , Exercise_Frequency , Sleep_Hours , Smoking_Habit , Alcohol_Consumption	Behavior patterns
Dietary Intake	Caloric_Intake , Protein_Intake , Carbohydrate_Intake , Fat_Intake , Preferred_Cuisine , Food_Aversions	Actual diet consumed
Recommendations	Recommended_Calories , Recommended_Protein , Recommended_Carbs , Recommended_Fats , Recommended_Meal_Plan	Personalized nutritional advice

In [608]:

```
missing_values = df.isnull().sum()
print(missing_values[missing_values > 0])
```

Chronic\_Disease 2043  
Allergies 3497  
Food\_Aversions 1225  
dtype: int64

In [609]:

```
desease_count = df['Chronic_Disease'].value_counts()
print(desease_count)
```

Chronic\_Disease  
Diabetes 1019  
Heart Disease 749  
Hypertension 693  
Obesity 496  
Name: count, dtype: int64

In [610]:

```
unique_dietary_habits = df['Dietary_Habits'].dropna().unique()

print("Unique Dietary Habits:")
for habit in sorted(unique_dietary_habits):
    print("-", habit)
```

Unique Dietary Habits:

- Keto
- Regular
- Vegan
- Vegetarian

In [611]:

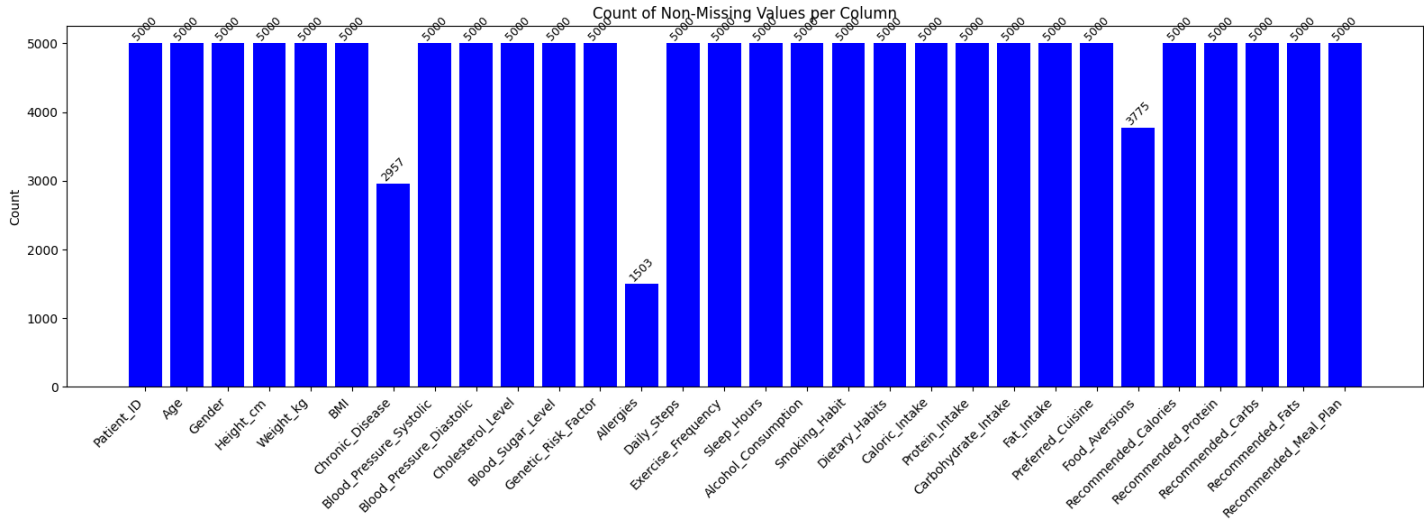
```
import pandas as pd
import matplotlib.pyplot as plt

non_missing_counts = df.notnull().sum()

plt.figure(figsize=(16, 6))
bars = plt.bar(non_missing_counts.index, non_missing_counts.values, color='blue')

for bar in bars:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, height, f'{int(height)}',
             ha='center', va='bottom', rotation=45, fontsize=9)

plt.xticks(rotation=45, ha='right')
plt.title("Count of Non-Missing Values per Column")
plt.ylabel("Count")
plt.tight_layout()
plt.show()
```



In [612]:

```
df['Recommended_Meal_Plan'].value_counts()
```

Out[612]:

count	
Recommended_Meal_Plan	
Low-Fat Diet	1313
High-Protein Diet	1255
Balanced Diet	1250
Low-Carb Diet	1182

dtype: int64

In [616]:

```
unique_dietary_habits = df['Dietary_Habits'].dropna().unique()
```

```
print("Unique Dietary Habits:")
for habit in sorted(unique_dietary_habits):
    print("-", habit)
```

Unique Dietary Habits:

- Keto
- Regular
- Vegan
- Vegetarian

## Preprocessing

In [615]:

```
# Creating Smart_Diet from logic
df['Smart_Diet'] = df.apply(
    lambda row: 'Low Calorie' if row['Caloric_Intake'] < 2000 else
                'High Protein' if row['Protein_Intake'] > 100 else
                'Balanced', axis=1)

# Map Smart_Diet to Recommended_Meal_Plan
df['Recommended_Meal_Plan'] = df['Smart_Diet'].map({
    'Low Calorie': 'Keto',
    'High Protein': 'Paleo',
    'Balanced': 'Mediterranean'
})
```

In [617]:

```
df['Recommended_Meal_Plan'].value_counts()
```

Out[617]:

		count
Recommended_Meal_Plan		
	Paleo	2173
	Keto	1732
	Mediterranean	1095

**dtype: int64**

In [618]:

```
# Separating columns
num_cols = df.select_dtypes(include=['int64', 'float64']).columns
cat_cols = df.select_dtypes(include='object').columns

# Fill numeric columns with median
df[num_cols] = df[num_cols].apply(lambda x: x.fillna(x.median()))

# Fill categorical columns with mode
df[cat_cols] = df[cat_cols].apply(lambda x: x.fillna(x.mode()[0]))

print("Missing values filled: median for numeric, mode for categorical")

# df.dropna(inplace=True)
```

Missing values filled: median for numeric, mode for categorical

In [619]:


```
import pandas as pd
import matplotlib.pyplot as plt
```

```
non_missing_counts = df.notnull().sum()
```

```
plt.figure(figsize=(16, 6))
bars = plt.bar(non_missing_counts.index, non_missing_counts.values, color='blue')

for bar in bars:
    height = bar.get_height()
    plt.text(bar.get_x() + bar.get_width()/2, height, f'{int(height)}',
             ha='center', va='bottom', rotation=45, fontsize=9)

plt.xticks(rotation=45, ha='right')
plt.title("Count of Non-Missing Values per Column")
plt.ylabel("Count")
plt.tight_layout()
plt.show()
```



In [670]:

```
# Encode categorical columns
cat_cols = df.select_dtypes(include='object').columns.difference(['Smart_Diet', 'Recommended_Meal_Plan'])
encoders = {}
for col in cat_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    encoders[col] = le
```

In [671]:

```
# Encode target labels BEFORE train_test_split
le_sd = LabelEncoder()
le_rmp = LabelEncoder()

df['Smart_Diet'] = le_sd.fit_transform(df['Smart_Diet'])
df['Recommended_Meal_Plan'] = le_rmp.fit_transform(df['Recommended_Meal_Plan'])
```

In [672]:

```
# Define features and targets
X = df.drop(['Smart_Diet', 'Recommended_Meal_Plan'], axis=1)
y = df[['Smart_Diet', 'Recommended_Meal_Plan']].copy()
```

In [673]:

```
# Check class distributions
print(f"Smart_Diet distribution: {Counter(y['Smart_Diet'])}")
print(f"Recommended_Meal_Plan distribution: {Counter(y['Recommended_Meal_Plan'])}")
```

```
Smart_Diet distribution: Counter({1: 2173, 2: 1732, 0: 1095})
Recommended_Meal_Plan distribution: Counter({2: 2173, 0: 1732, 1: 1095})
```

In [674]:

```
# Split dataset
X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.2, random_state=42, stratify=y
)
```

In [675]:

```
# Scale numerical features
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)
```

In [676]:

```
# Check for NaN values
if np.any(np.isnan(X_train_scaled)) or y_train.isna().any().any():
    raise ValueError("NaN values found in X_train_scaled or y_train")
```

In [677]:

```
# Before SMOTE Balance Smart_Diet
y_train_sd = y_train['Smart_Diet'].astype(int)
print(f"Before-SMOTE Smart_Diet distribution: {Counter(y_train_sd)}")
```

Before-SMOTE Smart\_Diet distribution: Counter({1: 1738, 2: 1386, 0: 876})

In [678]:

```
sd_counts = Counter(y_train_sd)
min_class_samples = min(sd_counts.values())
k_neighbors_safe = max(1, min(min_class_samples - 1, 5))

smote = SMOTE(k_neighbors=k_neighbors_safe, random_state=42)
X_train_bal_sd, y_train_bal_sd = smote.fit_resample(X_train_scaled, y_train_sd)
y_train_bal_sd = y_train_bal_sd.astype(int)
print(f"After-SMOTE Smart_Diet distribution: {Counter(y_train_bal_sd)}")
```

After-SMOTE Smart\_Diet distribution: Counter({1: 1738, 2: 1738, 0: 1738})

In [679]:

```
# Align Recommended_Meal_Plan with resampled Smart_Diet
original_indices = np.arange(len(X_train_scaled))
class_counts_before = Counter(y_train_sd)
class_counts_after = Counter(y_train_bal_sd)
sample_indices = []
for label in class_counts_after:
    original_mask = y_train_sd == label
    original_idx = original_indices[original_mask]
    n_original = class_counts_before[label]
    n_total = class_counts_after[label]
    sample_indices.extend(original_idx)
    if n_total > n_original:
        sample_indices.extend(np.repeat(original_idx[:1], n_total - n_original))
sample_indices = np.array(sample_indices)

y_train_bal = y_train.iloc[sample_indices].reset_index(drop=True)
y_train_bal['Smart_Diet'] = y_train_bal_sd
```

In [680]:

```
# SMOTE Balance Recommended_Meal_Plan
y_train_rmp = y_train_bal['Recommended_Meal_Plan'].astype(int)
if y_train_rmp.isna().any():
    raise ValueError("NaN values in y_train_bal['Recommended_Meal_Plan']")
print(f"Before-SMOTE Recommended_Meal_Plan distribution: {Counter(y_train_rmp)}")

rmp_counts = Counter(y_train_rmp)
min_class_samples = min(rmp_counts.values())
```



```

k_neighbors_safe = max(1, min(min_class_samples - 1, 5))

smote = SMOTE(k_neighbors=k_neighbors_safe, random_state=42)
X_train_bal, y_train_bal_rmp = smote.fit_resample(X_train_bal_sd, y_train_rmp)

y_train_bal_rmp = y_train_bal_rmp.astype(int)
print(f"After-SMOTE Recommended_Meal_Plan distribution: {Counter(y_train_bal_rmp)}")

```

Before-SMOTE Recommended\_Meal\_Plan distribution: Counter({2: 1738, 0: 1738, 1: 1738})  
 After-SMOTE Recommended\_Meal\_Plan distribution: Counter({2: 1738, 0: 1738, 1: 1738})

In [681]:

```

# Re-align Smart_Diet with resampled Recommended_Meal_Plan
original_indices = np.arange(len(y_train_bal))
class_counts_before = Counter(y_train_bal['Recommended_Meal_Plan'])
class_counts_after = Counter(y_train_bal_rmp)
sample_indices = []
for label in class_counts_after:
    original_mask = y_train_bal['Recommended_Meal_Plan'] == label
    original_idx = original_indices[original_mask]
    n_original = class_counts_before[label]
    n_total = class_counts_after[label]
    sample_indices.extend(original_idx)
    if n_total > n_original:
        sample_indices.extend(np.repeat(original_idx[:1], n_total - n_original))
sample_indices = np.array(sample_indices)

```

In [682]:

```

y_train_bal = y_train_bal.iloc[sample_indices].reset_index(drop=True)
y_train_bal['Recommended_Meal_Plan'] = y_train_bal_rmp
y_train_bal['Smart_Diet'] = y_train_bal['Smart_Diet'].astype(int)
print(f"y_train_bal dtypes: {y_train_bal.dtypes}")

```

```

y_train_bal dtypes: Smart_Diet          int64
Recommended_Meal_Plan    int64
dtype: object

```

In [683]:

```

# Train model
multi_model = MultiOutputClassifier(
    XGBClassifier(
        n_estimators=500,
        max_depth=7,
        learning_rate=0.05,
        subsample=0.8,
        colsample_bytree=0.8,
        random_state=42,
        eval_metric='mlogloss'
    )
)
multi_model.fit(X_train_bal, y_train_bal)

```

Out[683]:

```

► MultiOutputClassifier
  i ?

► estimator:
  XGBClassifier

► XGBClassifier

```

In [684]:

```

# Predict
y_pred = multi_model.predict(X_test_scaled)

```

In [685]:

```
# Evaluate
y_pred_df = pd.DataFrame(y_pred, columns=['Smart_Diet', 'Recommended_Meal_Plan'])
y_pred_df['Smart_Diet'] = le_sd.inverse_transform(y_pred_df['Smart_Diet'])
y_pred_df['Recommended_Meal_Plan'] = le_rmp.inverse_transform(y_pred_df['Recommended_Meal_Plan'])
y_test_decoded = y_test.copy()
y_test_decoded['Smart_Diet'] = le_sd.inverse_transform(y_test['Smart_Diet'])
y_test_decoded['Recommended_Meal_Plan'] = le_rmp.inverse_transform(y_test['Recommended_Meal_Plan'])

for col in ['Smart_Diet', 'Recommended_Meal_Plan']:
    acc = accuracy_score(y_test_decoded[col], y_pred_df[col])
    print(f" {col} Accuracy: {acc:.4f}")
    print(f"Classification Report for {col}:\n")
    print(classification_report(y_test_decoded[col], y_pred_df[col], zero_division=0))
    print("-" * 40)
```

Smart\_Diet Accuracy: 0.9990  
Classification Report for Smart\_Diet:

	precision	recall	f1-score	support
Balanced	1.00	1.00	1.00	219
High Protein	1.00	1.00	1.00	435
Low Calorie	1.00	1.00	1.00	346
accuracy			1.00	1000
macro avg	1.00	1.00	1.00	1000
weighted avg	1.00	1.00	1.00	1000

-----  
Recommended\_Meal\_Plan Accuracy: 0.4340  
Classification Report for Recommended\_Meal\_Plan:

	precision	recall	f1-score	support
Keto	0.39	0.54	0.45	346
Mediterranean	0.71	0.05	0.10	219
Paleo	0.47	0.54	0.50	435
accuracy			0.43	1000
macro avg	0.52	0.38	0.35	1000
weighted avg	0.49	0.43	0.40	1000

In [637]:

```
from sklearn.ensemble import RandomForestClassifier
import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Train on Recommended_Meal_Plan balanced data
rf = RandomForestClassifier(n_estimators=100, random_state=42)
rf.fit(X_train_bal, y_train_bal_rmp)

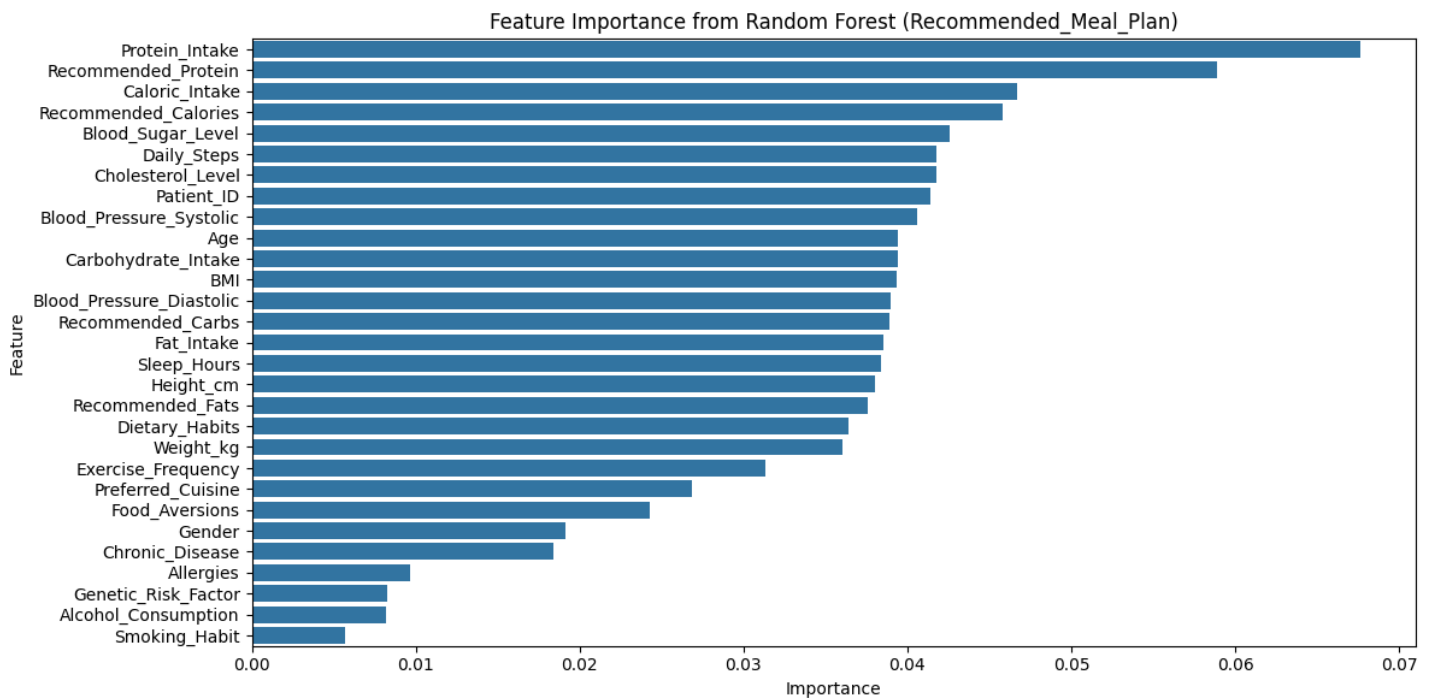
# Get feature importances
importances = rf.feature_importances_

# Assuming X_train_bal comes from X_train_scaled with these columns:
feature_names = df.drop(columns=['Smart_Diet', 'Recommended_Meal_Plan']).columns

# Create DataFrame
importance_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

# Plot
```

```
plt.figure(figsize=(12, 6))
sns.barplot(data=importance_df, x='Importance', y='Feature')
plt.title('Feature Importance from Random Forest (Recommended_Meal_Plan)')
plt.tight_layout()
plt.show()
```



In [638]:

```
# Imports
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt

from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier

from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, classification_report
from imblearn.over_sampling import SMOTE

# Encode categorical features (excluding targets)
cat_cols = df.select_dtypes(include='object').columns.difference(['Smart_Diet', 'Recommended_Meal_Plan'])
encoders = {}
for col in cat_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    encoders[col] = le

# Encode the target column (Recommended_Meal_Plan)
le_rmp = LabelEncoder()
df['Recommended_Meal_Plan'] = le_rmp.fit_transform(df['Recommended_Meal_Plan'])

# Feature and target separation
X = df.drop(columns=['Smart_Diet', 'Recommended_Meal_Plan'])
y = df['Recommended_Meal_Plan']

# Split the dataset before applying SMOTE
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)

# Apply SMOTE on training data only
smote = SMOTE(random_state=42)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)
```

```

# Scale the data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_bal)
X_test_scaled = scaler.transform(X_test)

# Define classifiers
models = {
    "Random Forest": RandomForestClassifier(n_estimators=100, random_state=42),
    "XGBoost": XGBClassifier(eval_metric='mlogloss', random_state=42),
    "Logistic Regression": LogisticRegression(max_iter=1000, solver='lbfgs', random_state=42)
}

# Train and evaluate each model
for name, model in models.items():
    print(f"\n{name}")

    # Train
    model.fit(X_train_scaled, y_train_bal)

    # Predict
    y_pred = model.predict(X_test_scaled)

    # Evaluation
    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred, average='weighted')
    print("Accuracy:", (acc*100))
    print("F1 Score:", f1)

    print("Classification Report:\n", classification_report(
        y_test, y_pred, target_names=le_rmp.classes_.astype(str)
    ))

    # Confusion Matrix
    cm = confusion_matrix(y_test, y_pred)
    # plt.figure(figsize=(6, 4))
    # sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    # plt.title(f'Confusion Matrix - {name}')
    # plt.xlabel('Predicted')
    # plt.ylabel('Actual')
    # plt.tight_layout()
    # plt.show()
    print("--"*40)

```

Random Forest  
Accuracy: 100.0  
F1 Score: 1.0  
Classification Report:

	precision	recall	f1-score	support
0	1.00	1.00	1.00	346
1	1.00	1.00	1.00	219
2	1.00	1.00	1.00	435
accuracy			1.00	1000
macro avg	1.00	1.00	1.00	1000
weighted avg	1.00	1.00	1.00	1000

-----

XGBoost  
Accuracy: 99.7  
F1 Score: 0.9969981513990142  
Classification Report:

	precision	recall	f1-score	support
0	0.99	1.00	1.00	346
1	1.00	0.99	1.00	219
2	1.00	1.00	1.00	435
accuracy			1.00	1000

macro avg	1.00	1.00	1.00	1000
weighted avg	1.00	1.00	1.00	1000

-----

Logistic Regression

Accuracy: 97.0

F1 Score: 0.9700353879941769

Classification Report:

	precision	recall	f1-score	support
0	0.95	0.98	0.97	346
1	0.96	0.96	0.96	219
2	0.99	0.97	0.98	435
accuracy			0.97	1000
macro avg	0.97	0.97	0.97	1000
weighted avg	0.97	0.97	0.97	1000

-----

In [639]:

```
from sklearn.model_selection import cross_val_score

for name, model in models.items():
    print(f"\n{name} (Cross-Validation Accuracy)")
    scores = cross_val_score(model, X, y, cv=5, scoring='accuracy')
    print("CV Mean Accuracy:", scores.mean())
    print("CV Std Dev:", scores.std())
```

Random Forest (Cross-Validation Accuracy)

CV Mean Accuracy: 0.9996

CV Std Dev: 0.000489897948556636

XGBoost (Cross-Validation Accuracy)

CV Mean Accuracy: 0.9982000000000001

CV Std Dev: 0.000979795897113272

Logistic Regression (Cross-Validation Accuracy)

CV Mean Accuracy: 0.9456

CV Std Dev: 0.00320000000000000028

In [661]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, classification_report
from imblearn.over_sampling import SMOTE
from sklearn.feature_selection import SelectFromModel
from sklearn.impute import SimpleImputer

# Verify columns
print("Columns in DataFrame:", df.columns.tolist())

# Add noise to reduce deterministic mapping
np.random.seed(42)
df['Carbohydrate_Intake'] += np.random.normal(0, 20, df.shape[0])
df['Fat_Intake'] += np.random.normal(0, 15, df.shape[0])

# Redefine target with adjusted thresholds
df['Smart_Diet'] = df.apply(
    lambda row: 'Low Carb' if row['Carbohydrate_Intake'] < 180 else
    'High Fat' if row['Fat_Intake'] > 90 else
```

```

        'Balanced', axis=1)
df['Recommended_Meal_Plan'] = df['Smart_Diet'].map({
    'Low Carb': 'Keto',
    'High Fat': 'Paleo',
    'Balanced': 'Mediterranean'
})
df = df.drop(columns=['Smart_Diet'])

# Drop non-predictive columns
columns_to_drop = [col for col in ['Patient_ID', 'BMI'] if col in df.columns]
df = df.drop(columns=columns_to_drop)

# Impute missing values before encoding
imputer = SimpleImputer(strategy='constant', fill_value='Unknown')
for col in ['Chronic_Disease', 'Allergies', 'Food_Aversions']:
    if col in df.columns and df[col].dtype == 'object':
        df[col] = imputer.fit_transform(df[[col]]).ravel()

# Check for highly correlated features
corr_matrix = df.select_dtypes(include=['int64', 'float64']).corr()
high_corr = [(i, j) for i in corr_matrix for j in corr_matrix if corr_matrix.loc[i, j] >
0.8 and i < j]
print("Highly correlated features:", high_corr)

# Encode categorical features
cat_cols = df.select_dtypes(include='object').columns.difference(['Recommended_Meal_Plan'])
encoders = {}
for col in cat_cols:
    if df[col].dtype == 'object':
        le = LabelEncoder()
        df[col] = le.fit_transform(df[col])
        encoders[col] = le

# Encode target column
le_rmp = LabelEncoder()
df['Recommended_Meal_Plan'] = le_rmp.fit_transform(df['Recommended_Meal_Plan'])
print("Class mappings:", dict(zip(range(len(le_rmp.classes_)), le_rmp.classes_)))

# Split features and target
X = df.drop(columns=['Recommended_Meal_Plan'])
y = df['Recommended_Meal_Plan']

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)

# Step 8: Print class distribution
train_class_counts = pd.Series(y_train).value_counts()
print("Training set class distribution:", train_class_counts.to_dict())

# Apply SMOTE with corrected sampling strategy
smote = SMOTE(
    sampling_strategy={0: 1646, 1: 1646, 2: 1646}, # Balance classes
    random_state=42,
    k_neighbors=5
)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)

# Scale data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_bal)
X_test_scaled = scaler.transform(X_test)

# Feature selection using Random Forest
rf_for_selection = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)
rf_for_selection.fit(X_train_scaled, y_train_bal)
selector = SelectFromModel(rf_for_selection, prefit=True, threshold="0.5*mean")
X_train_selected = selector.transform(X_train_scaled)
X_test_selected = selector.transform(X_test_scaled)
print("Selected features shape:", X_train_selected.shape)
selected_features = X.columns[selector.get_support()].tolist()

```

```

print("Selected features:", selected_features)
print("Feature importances:", dict(zip(X.columns, rf_for_selection.feature_importances_))
)

# Define tuned classifiers
models = {
    "XGBoost": XGBClassifier(
        max_depth=3, reg_lambda=12.0, reg_alpha=10.0, subsample=0.7, colsample_bytree=0.
7,
        eval_metric='mlogloss', random_state=42
    ),
    "Logistic Regression": GridSearchCV(
        LogisticRegression(max_iter=2000, random_state=42),
        param_grid={
            'C': [0.01, 0.1, 0.5, 1.0, 2.0, 5.0, 10.0, 20.0],
            'solver': ['lbfgs', 'saga'],
            'class_weight': ['balanced']
        },
        cv=5,
        scoring='accuracy',
        n_jobs=-1
    )
}

# Train and evaluate models
for name, model in models.items():
    print(f"\n{name}")

    model.fit(X_train_selected, y_train_bal)
    y_pred = model.predict(X_test_selected)

    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred, average='weighted')

    print("Accuracy:", round(acc * 100, 2))
    print("F1 Score:", round(f1, 4))
    print("Classification Report:\n", classification_report(y_test, y_pred, target_names
=le_rmp.classes_))

    cv_scores = cross_val_score(model, X_train_selected, y_train_bal, cv=5, scoring='acc
uracy')
    print("Cross-Validation Accuracy (mean ± std):", f"{cv_scores.mean() * 100:.2f}% ± {
cv_scores.std() * 100:.2f}%")

    cm = confusion_matrix(y_test, y_pred)
    # Uncomment to visualize confusion matrix
    # plt.figure(figsize=(6, 4))
    # sns.heatmap(cm, annot=True, fmt='d', cmap='Blues')
    # plt.title(f'Confusion Matrix - {name}')
    # plt.xlabel('Predicted')
    # plt.ylabel('Actual')
    # plt.tight_layout()
    # plt.show()

    print("--" * 40)

```

Columns in DataFrame: ['Age', 'Gender', 'Height\_cm', 'Weight\_kg', 'Chronic\_Disease', 'Blood\_Pressure\_Systolic', 'Blood\_Pressure\_Diastolic', 'Cholesterol\_Level', 'Blood\_Sugar\_Level', 'Genetic\_Risk\_Factor', 'Allergies', 'Daily\_Steps', 'Exercise\_Frequency', 'Sleep\_Hours', 'Alcohol\_Consumption', 'Smoking\_Habit', 'Dietary\_Habits', 'Caloric\_Intake', 'Protein\_Intake', 'Carbohydrate\_Intake', 'Fat\_Intake', 'Preferred\_Cuisine', 'Food\_Aversions', 'Recommended\_Calories', 'Recommended\_Protein', 'Recommended\_Carbs', 'Recommended\_Fats', 'Recommended\_Meal\_Plan', 'Health\_Score', 'Risk\_Score']

Highly correlated features: [('Caloric\_Intake', 'Recommended\_Calories'), ('Protein\_Intake', 'Recommended\_Protein')]

Class mappings: {0: 'Keto', 1: 'Mediterranean', 2: 'Paleo'}

Training set class distribution: {0: 1442, 1: 1335, 2: 1223}

Selected features shape: (4938, 3)

Selected features: ['Carbohydrate\_Intake', 'Fat\_Intake', 'Recommended\_Carbs']

Feature importances: {'Age': np.float64(0.006472920471064673), 'Gender': np.float64(0.0015056720390958906), 'Height\_cm': np.float64(0.006134038422741217), 'Weight\_kg': np.float64(0.006186817682737418), 'Chronic\_Disease': np.float64(0.0021266824919182747), 'Blood\_Pres

```
sure_Systolic': np.float64(0.0060899057180364225), 'Blood_Pressure_Diastolic': np.float64(0.005907394278162885), 'Cholesterol_Level': np.float64(0.006756306736900678), 'Blood_Sugar_Level': np.float64(0.006320871417388234), 'Genetic_Risk_Factor': np.float64(0.0009932782699434515), 'Allergies': np.float64(0.001295282659855451), 'Daily_Steps': np.float64(0.008152339381582602), 'Exercise_Frequency': np.float64(0.0033095469870122298), 'Sleep_Hours': np.float64(0.006816640468274531), 'Alcohol_Consumption': np.float64(0.00113732414562591), 'Smoking_Habit': np.float64(0.0007020585747328368), 'Dietary_Habits': np.float64(0.016952946844144708), 'Caloric_Intake': np.float64(0.007020656526342036), 'Protein_Intake': np.float64(0.005872363974178726), 'Carbohydrate_Intake': np.float64(0.49053369762880344), 'Fat_Intake': np.float64(0.3540697713292837), 'Preferred_Cuisine': np.float64(0.002147922030460854), 'Food_Aversions': np.float64(0.0014050186670294881), 'Recommended_Calories': np.float64(0.006923747369427264), 'Recommended_Protein': np.float64(0.00663885160970012), 'Recommended_Carbs': np.float64(0.03827981802914785), 'Recommended_Fats': np.float64(0.014598331375268224), 'Health_Score': np.float64(0.0009074470308711066), 'Risk_Score': np.float64(0.0)}
```

XGBoost

Accuracy: 99.7

F1 Score: 0.997

Classification Report:

	precision	recall	f1-score	support
Keto	1.00	1.00	1.00	361
Mediterranean	1.00	0.99	1.00	333
Paleo	0.99	1.00	1.00	306
accuracy			1.00	1000
macro avg	1.00	1.00	1.00	1000
weighted avg	1.00	1.00	1.00	1000

Cross-Validation Accuracy (mean ± std): 99.76% ± 0.26%

Logistic Regression

Accuracy: 98.7

F1 Score: 0.987

Classification Report:

	precision	recall	f1-score	support
Keto	0.98	0.99	0.98	361
Mediterranean	1.00	0.98	0.99	333
Paleo	0.99	0.98	0.98	306
accuracy			0.99	1000
macro avg	0.99	0.99	0.99	1000
weighted avg	0.99	0.99	0.99	1000

Cross-Validation Accuracy (mean ± std): 98.46% ± 0.71%

In [641]:

```
print(df.columns.tolist())
```

```
['Age', 'Gender', 'Height_cm', 'Weight_kg', 'Chronic_Disease', 'Blood_Pressure_Systolic', 'Blood_Pressure_Diastolic', 'Cholesterol_Level', 'Blood_Sugar_Level', 'Genetic_Risk_Factor', 'Allergies', 'Daily_Steps', 'Exercise_Frequency', 'Sleep_Hours', 'Alcohol_Consumption', 'Smoking_Habit', 'Dietary_Habits', 'Caloric_Intake', 'Protein_Intake', 'Carbohydrate_Intake', 'Fat_Intake', 'Preferred_Cuisine', 'Food_Aversions', 'Recommended_Calories', 'Recommended_Protein', 'Recommended_Carbs', 'Recommended_Fats', 'Recommended_Meal_Plan']
```

In [ ]:

```
df.info()
```

In [662]:

```
from sklearn.preprocessing import PolynomialFeatures
poly = PolynomialFeatures(degree=2, interaction_only=True, include_bias=False)
X_poly = poly.fit_transform(X)
```



```
In [663]:
```

```
from sklearn.feature_selection import SelectKBest, f_classif
selector = SelectKBest(score_func=f_classif, k=20) # pick top 20 features
X_train_selected = selector.fit_transform(X_train_scaled, y_train_bal)
X_test_selected = selector.transform(X_test_scaled)
```

```
In [664]:
```

```
LogisticRegression(penalty='l1', solver='saga')
```

```
Out[664]:
```

```
▼ LogisticRegression i ?
LogisticRegression(penalty='l1', solver='saga')
```

```
In [665]:
```

```
# Train Logistic Regression GridSearchCV
logreg_model = models["Logistic Regression"]
logreg_model.fit(X_train_selected, y_train_bal)

# Extract the best model
best_logistic_regression = logreg_model.best_estimator_
```

```
In [646]:
```

```
xgb_model = XGBClassifier(
    max_depth=2, reg_lambda=6.0, reg_alpha=5.0, eval_metric='mlogloss', random_state=42
)
xgb_model.fit(X_train_selected, y_train_bal)
best_xgboost_model = xgb_model
```

## APPLYING ENSEMBLE

```
In [647]:
```

```
from sklearn.ensemble import VotingClassifier

ensemble = VotingClassifier(estimators=[
    ('lr', best_logistic_regression),
    ('xgb', best_xgboost_model)
], voting='soft')

ensemble.fit(X_train_selected, y_train_bal)
y_pred = ensemble.predict(X_test_selected)
```

```
In [648]:
```

```
from sklearn.metrics import accuracy_score, f1_score, classification_report

acc = accuracy_score(y_test, y_pred)
f1 = f1_score(y_test, y_pred, average='weighted')

print("Ensemble Accuracy:", acc * 100)
print("Ensemble F1 Score:", f1)
print("Classification Report:\n", classification_report(
    y_test, y_pred, target_names=le_rmp.classes_.astype(str)
))
```

```
Ensemble Accuracy: 99.3
```

```
Ensemble F1 Score: 0.9930006450254993
```

```
Classification Report:
```

	precision	recall	f1-score	support
Keto	0.99	1.00	0.99	271
Mediterranean	1.00	0.99	0.99	393
Paleo	0.99	1.00	1.00	336

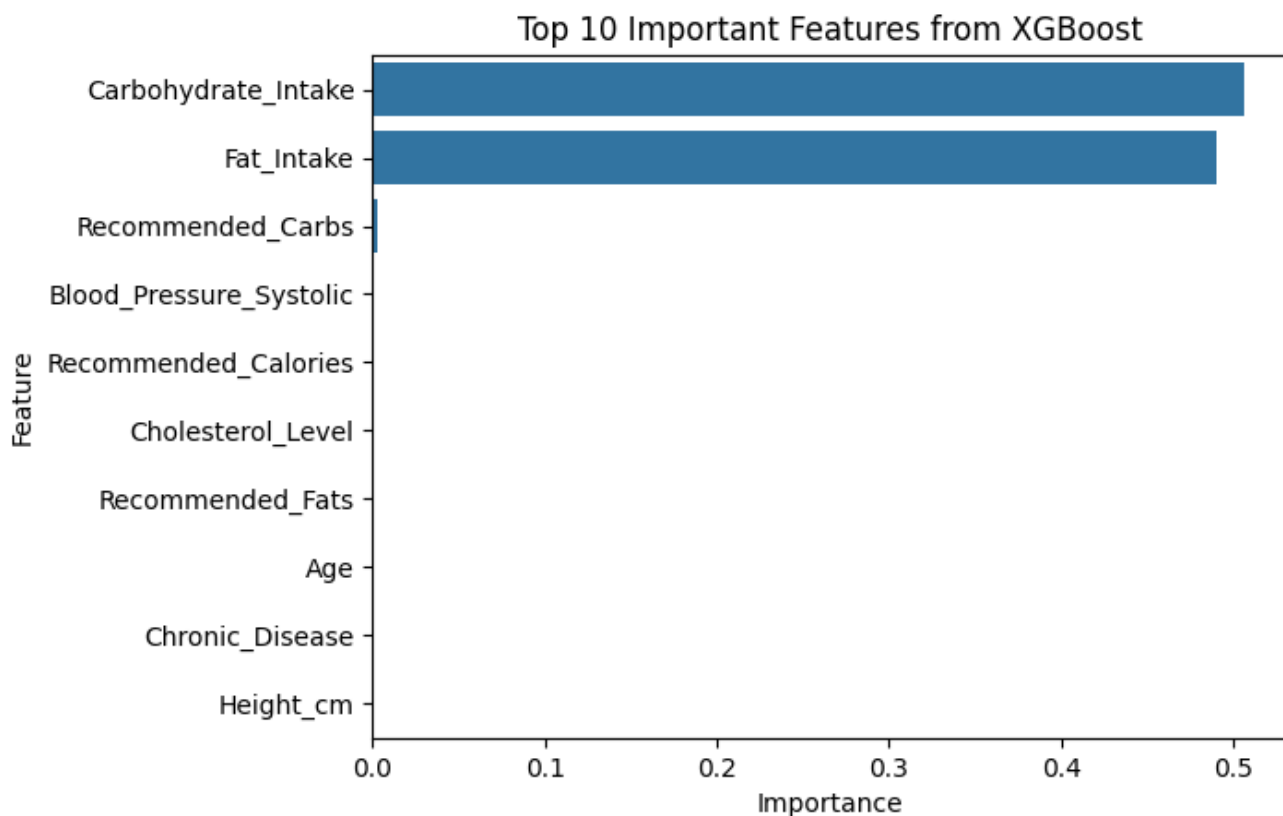
accuracy			0.99	1000
macro avg	0.99	0.99	0.99	1000
weighted avg	0.99	0.99	0.99	1000

In [649]:

```
importances = best_xgboost_model.feature_importances_
selected_features = selector.get_support(indices=True)
feature_names = X.columns[selected_features]

importance_df = pd.DataFrame({
    'Feature': feature_names,
    'Importance': importances
}).sort_values(by='Importance', ascending=False)

sns.barplot(data=importance_df.head(10), x='Importance', y='Feature')
plt.title("Top 10 Important Features from XGBoost")
plt.show()
```



In [650]:

```
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import accuracy_score, f1_score, confusion_matrix, classification_report
from imblearn.over_sampling import SMOTE
from sklearn.feature_selection import SelectFromModel
from sklearn.impute import SimpleImputer

# Verify columns
print("Columns in DataFrame:", df.columns.tolist())

# Add noise to reduce deterministic mapping
np.random.seed(42)
df['Carbohydrate_Intake'] += np.random.normal(0, 25, df.shape[0])
```

```

df['Fat_Intake'] += np.random.normal(0, 20, df.shape[0])

# Redefine target with adjusted thresholds
df['Smart_Diet'] = df.apply(
    lambda row: 'Low Carb' if row['Carbohydrate_Intake'] < 180 else
                'High Fat' if row['Fat_Intake'] > 90 else
                'Balanced', axis=1)
df['Recommended_Meal_Plan'] = df['Smart_Diet'].map({
    'Low Carb': 'Keto',
    'High Fat': 'Paleo',
    'Balanced': 'Mediterranean'
})
df = df.drop(columns=['Smart_Diet'])

# Drop non-predictive columns
columns_to_drop = [col for col in ['Patient_ID', 'BMI'] if col in df.columns]
df = df.drop(columns=columns_to_drop)

# Impute missing values before encoding
imputer = SimpleImputer(strategy='constant', fill_value='Unknown')
for col in ['Chronic_Disease', 'Allergies', 'Food_Aversions']:
    if col in df.columns and df[col].dtype == 'object':
        df[col] = imputer.fit_transform(df[[col]]).ravel()

# Check for highly correlated features
corr_matrix = df.select_dtypes(include=['int64', 'float64']).corr()
high_corr = [(i, j) for i in corr_matrix for j in corr_matrix if corr_matrix.loc[i, j] >
0.8 and i < j]
print("Highly correlated features:", high_corr)

# Encode categorical features
cat_cols = df.select_dtypes(include='object').columns.difference(['Recommended_Meal_Plan'])
encoders = {}
for col in cat_cols:
    if df[col].dtype == 'object':
        le = LabelEncoder()
        df[col] = le.fit_transform(df[col])
        encoders[col] = le

# Encode target column
le_rmp = LabelEncoder()
df['Recommended_Meal_Plan'] = le_rmp.fit_transform(df['Recommended_Meal_Plan'])
print("Class mappings:", dict(zip(range(len(le_rmp.classes_)), le_rmp.classes_)))

# Split features and target
X = df.drop(columns=['Recommended_Meal_Plan'])
y = df['Recommended_Meal_Plan']

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)

# Print class distribution
train_class_counts = pd.Series(y_train).value_counts()
print("Training set class distribution:", train_class_counts.to_dict())

# Apply SMOTE with corrected sampling strategy
smote = SMOTE(
    sampling_strategy={0: 1646, 1: 1646, 2: 1646}, # Balance classes
    random_state=42,
    k_neighbors=5
)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)

# Scale data
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_bal)
X_test_scaled = scaler.transform(X_test)

# Feature selection using Random Forest
rf_for_selection = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)

```

```

rf_for_selection.fit(X_train_scaled, y_train_bal)
selector = SelectFromModel(rf_for_selection, prefit=True, threshold="0.5*mean")
X_train_selected = selector.transform(X_train_scaled)
X_test_selected = selector.transform(X_test_scaled)
print("Selected features shape:", X_train_selected.shape)
selected_features = X.columns[selector.get_support()].tolist()
print("Selected features:", selected_features)
print("Feature importances:", dict(zip(X.columns, rf_for_selection.feature_importances_))
)

# Define Logistic Regression model
logistic_model = GridSearchCV(
    LogisticRegression(max_iter=2000, random_state=42),
    param_grid={
        'C': [0.01, 0.1, 0.5, 1.0, 2.0, 5.0, 10.0, 20.0, 50.0],
        'solver': ['lbfgs', 'saga'],
        'class_weight': ['balanced']
    },
    cv=5,
    scoring='accuracy',
    n_jobs=-1
)

# Train and evaluate Logistic Regression
print("\nLogistic Regression")

logistic_model.fit(X_train_selected, y_train_bal)
y_pred_logistic = logistic_model.predict(X_test_selected)

acc_logistic = accuracy_score(y_test, y_pred_logistic)
f1_logistic = f1_score(y_test, y_pred_logistic, average='weighted')

print("Accuracy:", round(acc_logistic * 100, 2))
print("F1 Score:", round(f1_logistic, 4))
print("Classification Report:\n", classification_report(y_test, y_pred_logistic, target_names=le_rmp.classes_))

cv_scores_logistic = cross_val_score(logistic_model, X_train_selected, y_train_bal, cv=5,
    , scoring='accuracy')
print("Cross-Validation Accuracy (mean ± std):", f"{cv_scores_logistic.mean() * 100:.2f}% ± {cv_scores_logistic.std() * 100:.2f}%")

# Confusion matrix (optional visualization)
# cm_logistic = confusion_matrix(y_test, y_pred_logistic)
# plt.figure(figsize=(6, 4))
# sns.heatmap(cm_logistic, annot=True, fmt='d', cmap='Blues')
# plt.title('Confusion Matrix - Logistic Regression')
# plt.xlabel('Predicted')
# plt.ylabel('Actual')
# plt.tight_layout()
# plt.show()

print("---" * 40)

```

Columns in DataFrame: ['Age', 'Gender', 'Height\_cm', 'Weight\_kg', 'Chronic\_Disease', 'Blood\_Pressure\_Systolic', 'Blood\_Pressure\_Diastolic', 'Cholesterol\_Level', 'Blood\_Sugar\_Level', 'Genetic\_Risk\_Factor', 'Allergies', 'Daily\_Steps', 'Exercise\_Frequency', 'Sleep\_Hours', 'Alcohol\_Consumption', 'Smoking\_Habit', 'Dietary\_Habits', 'Caloric\_Intake', 'Protein\_Intake', 'Carbohydrate\_Intake', 'Fat\_Intake', 'Preferred\_Cuisine', 'Food\_Aversions', 'Recommended\_Calories', 'Recommended\_Protein', 'Recommended\_Carbs', 'Recommended\_Fats', 'Recommended\_Meal\_Plan']

Highly correlated features: [('Caloric\_Intake', 'Recommended\_Calories'), ('Protein\_Intake', 'Recommended\_Protein'), ('Carbohydrate\_Intake', 'Recommended\_Carbs')]

Class mappings: {0: 'Keto', 1: 'Mediterranean', 2: 'Paleo'}

Training set class distribution: {1: 1561, 2: 1338, 0: 1101}

Selected features shape: (4938, 4)

Selected features: ['Carbohydrate\_Intake', 'Fat\_Intake', 'Recommended\_Carbs', 'Recommended\_Fats']

Feature importances: {'Age': np.float64(0.0033317870577560423), 'Gender': np.float64(0.0010690640246683368), 'Height\_cm': np.float64(0.0038308378903513956), 'Weight\_kg': np.float64(0.0036213265621357106), 'Chronic\_Disease': np.float64(0.0014716712717129576), 'Blood\_Pressure\_Systolic': np.float64(0.003449709772963038), 'Blood\_Pressure\_Diastolic': np.float

```
64(0.0035102481767330493), 'Cholesterol_Level': np.float64(0.004117790072079343), 'Blood_Sugar_Level': np.float64(0.0036430953213339693), 'Genetic_Risk_Factor': np.float64(0.0005994480932218619), 'Allergies': np.float64(0.0007493635229322185), 'Daily_Steps': np.float64(0.004671598568497702), 'Exercise_Frequency': np.float64(0.0020691623501544413), 'Sleep_Hours': np.float64(0.0037825493985234253), 'Alcohol_Consumption': np.float64(0.0003680222958298461), 'Smoking_Habit': np.float64(0.0004600590872574526), 'Dietary_Habits': np.float64(0.0010962490231366424), 'Caloric_Intake': np.float64(0.004039968726301402), 'Protein_Intake': np.float64(0.004111014956678106), 'Carbohydrate_Intake': np.float64(0.4189945453060622), 'Fat_Intake': np.float64(0.3220070193850096), 'Preferred_Cuisine': np.float64(0.0016386837182538207), 'Food_Aversions': np.float64(0.0011755135054509435), 'Recommended_Calories': np.float64(0.004579867407605912), 'Recommended_Protein': np.float64(0.004032400842425556), 'Recommended_Carbs': np.float64(0.12515073083006217), 'Recommended_Fats': np.float64(0.07242827283286286)}
```

Logistic Regression

Accuracy: 98.2

F1 Score: 0.982

Classification Report:

	precision	recall	f1-score	support
Keto	0.96	0.97	0.97	275
Mediterranean	0.99	0.98	0.99	390
Paleo	0.99	0.99	0.99	335
accuracy			0.98	1000
macro avg	0.98	0.98	0.98	1000
weighted avg	0.98	0.98	0.98	1000

Cross-Validation Accuracy (mean ± std): 98.52% ± 0.80%

In [651]:

```
!pip install catboost
```

```
Requirement already satisfied: catboost in /usr/local/lib/python3.11/dist-packages (1.2.8)
Requirement already satisfied: graphviz in /usr/local/lib/python3.11/dist-packages (from catboost) (0.20.3)
Requirement already satisfied: matplotlib in /usr/local/lib/python3.11/dist-packages (from catboost) (3.10.0)
Requirement already satisfied: numpy<3.0,>=1.16.0 in /usr/local/lib/python3.11/dist-packages (from catboost) (2.0.2)
Requirement already satisfied: pandas>=0.24 in /usr/local/lib/python3.11/dist-packages (from catboost) (2.2.2)
Requirement already satisfied: scipy in /usr/local/lib/python3.11/dist-packages (from catboost) (1.15.3)
Requirement already satisfied: plotly in /usr/local/lib/python3.11/dist-packages (from catboost) (5.24.1)
Requirement already satisfied: six in /usr/local/lib/python3.11/dist-packages (from catboost) (1.17.0)
Requirement already satisfied: python-dateutil>=2.8.2 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.24->catboost) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: tzdata>=2022.7 in /usr/local/lib/python3.11/dist-packages (from pandas>=0.24->catboost) (2025.2)
Requirement already satisfied: contourpy>=1.0.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (1.3.2)
Requirement already satisfied: cyclor>=0.10 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (0.12.1)
Requirement already satisfied: fonttools>=4.22.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (4.58.1)
Requirement already satisfied: kiwisolver>=1.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (1.4.8)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (24.2)
Requirement already satisfied: pillow>=8 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (11.2.1)
Requirement already satisfied: pyparsing>=2.3.1 in /usr/local/lib/python3.11/dist-packages (from matplotlib->catboost) (3.2.3)
Requirement already satisfied: tenacity>=6.2.0 in /usr/local/lib/python3.11/dist-packages
```

```
(from plotly->catboost) (9.1.2)
```

```
In [651]:
```

```
In [693]:
```

```
# ----- Imports -----
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from collections import Counter

from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import accuracy_score, f1_score, classification_report, confusion_matrix

from sklearn.linear_model import LogisticRegression
from sklearn.ensemble import RandomForestClassifier
from xgboost import XGBClassifier
from sklearn.svm import SVC
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier

from imblearn.over_sampling import SMOTE

# ----- Preprocessing -----

# Add Gaussian noise to reduce overfit
np.random.seed(42)
df['Carbohydrate_Intake'] += np.random.normal(0, 20, df.shape[0])
df['Fat_Intake'] += np.random.normal(0, 15, df.shape[0])

# Create target label
df['Smart_Diet'] = df.apply(
    lambda row: 'Low Carb' if row['Carbohydrate_Intake'] < 180 else
                'High Fat' if row['Fat_Intake'] > 90 else
                'Balanced', axis=1)
df['Recommended_Meal_Plan'] = df['Smart_Diet'].map({
    'Low Carb': 'Keto',
    'High Fat': 'Paleo',
    'Balanced': 'Mediterranean'
})
df.drop(columns=['Smart_Diet'], inplace=True)

# Drop unnecessary columns
df.drop(columns=[col for col in ['Patient_ID', 'BMI'] if col in df.columns], inplace=True)

# Handle missing values
imputer = SimpleImputer(strategy='constant', fill_value='Unknown')
for col in ['Chronic_Disease', 'Allergies', 'Food_Aversions']:
    if col in df.columns and df[col].dtype == 'object':
        df[col] = imputer.fit_transform(df[[col]].ravel())

# Encode categorical features
cat_cols = df.select_dtypes(include='object').columns.difference(['Recommended_Meal_Plan'])
encoders = {}
for col in cat_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    encoders[col] = le

# Encode target
```

```

le_rmp = LabelEncoder()
df['Recommended_Meal_Plan'] = le_rmp.fit_transform(df['Recommended_Meal_Plan'])

# Split features and target
X = df.drop(columns=['Recommended_Meal_Plan'])
y = df['Recommended_Meal_Plan']

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)

# Print original class distribution
print("Original training class distribution:", Counter(y_train))

# Get max count among classes
max_count = max(Counter(y_train).values())

# Create sampling strategy to oversample minority classes to match max_count
sampling_strategy = {cls: max_count for cls in np.unique(y_train)}

# Apply SMOTE
smote = SMOTE(sampling_strategy=sampling_strategy, random_state=42)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)

print("After SMOTE:", Counter(y_train_bal))

# Add noise to X_train_bal
X_train_bal += np.random.normal(0, 0.01, X_train_bal.shape)

# Standard scaling
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_bal)
X_test_scaled = scaler.transform(X_test)

# Feature selection using RF
rf = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)
rf.fit(X_train_scaled, y_train_bal)
selector = SelectFromModel(rf, threshold="1.0*mean", prefit=True)
X_train_sel = selector.transform(X_train_scaled)
X_test_sel = selector.transform(X_test_scaled)
selected_features = X.columns[selector.get_support()].tolist()
print("Selected Features:", selected_features)

# ----- Model Setup -----
models = {
    "Random Forest": RandomForestClassifier(
        n_estimators=50, max_depth=4, min_samples_split=100,
        min_samples_leaf=40, max_features='sqrt', random_state=42
    ),
    "XGBoost": XGBClassifier(
        max_depth=3, reg_lambda=25.0, reg_alpha=15.0, subsample=0.6,
        colsample_bytree=0.6, learning_rate=0.07, eval_metric='mlogloss',
        use_label_encoder=False, random_state=42
    ),
    "Logistic Regression": GridSearchCV(
        LogisticRegression(
            max_iter=1000, penalty='elasticnet', solver='saga',
            l1_ratio=0.5, random_state=42
        ),
        param_grid={'C': [0.1, 0.5, 1.0], 'class_weight': ['balanced']},
        cv=5, scoring='accuracy', n_jobs=-1
    ),
    "SVM": SVC(
        kernel='rbf', C=1.0, gamma='scale', class_weight='balanced', probability=True, random_state=42
    ),
    "LightGBM": LGBMClassifier(
        num_leaves=20,
        max_depth=4,
        learning_rate=0.05,
        n_estimators=100,

```

```

        subsample=0.7,
        colsample_bytree=0.7,
        reg_lambda=5.0,
        reg_alpha=3.0,
        random_state=42,
        verbose=-1
    ),
    "CatBoost": CatBoostClassifier(
        iterations=100, depth=3, learning_rate=0.05,
        l2_leaf_reg=10, verbose=0, random_state=42
    )
}

# ----- Evaluation -----
for name, model in models.items():
    print(f"\nModel: {name}")
    model.fit(X_train_sel, y_train_bal)
    y_pred = model.predict(X_test_sel)

    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred, average='weighted')

    print("Accuracy:", round(acc * 100, 2), "%")
    print("F1 Score:", round(f1, 4))
    print("Classification Report:\n", classification_report(y_test, y_pred, target_names
=le_rmp.classes_))

    cv = cross_val_score(model, X_train_sel, y_train_bal, cv=5, scoring='accuracy')
    print(f"CV Accuracy: {cv.mean() * 100:.2f}% ± {cv.std() * 100:.2f}%")
    print("-" * 80)

```

Original training class distribution: Counter({0: 1654, 1: 1232, 2: 1114})

After SMOTE: Counter({0: 1654, 2: 1654, 1: 1654})

Selected Features: ['Carbohydrate\_Intake', 'Fat\_Intake']

Model: Random Forest

Accuracy: 100.0 %

F1 Score: 1.0

Classification Report:

	precision	recall	f1-score	support
Keto	1.00	1.00	1.00	413
Mediterranean	1.00	1.00	1.00	308
Paleo	1.00	1.00	1.00	279
accuracy			1.00	1000
macro avg	1.00	1.00	1.00	1000
weighted avg	1.00	1.00	1.00	1000

CV Accuracy: 99.94% ± 0.08%

Model: XGBoost

Accuracy: 99.6 %

F1 Score: 0.996

Classification Report:

	precision	recall	f1-score	support
Keto	1.00	0.99	1.00	413
Mediterranean	0.99	1.00	1.00	308
Paleo	1.00	1.00	1.00	279
accuracy			1.00	1000
macro avg	1.00	1.00	1.00	1000
weighted avg	1.00	1.00	1.00	1000

CV Accuracy: 99.74% ± 0.15%

Model: Logistic Regression

Accuracy: 97.8 %

F1 Score: 0.978



# Classification Report:

	precision	recall	f1-score	support
Keto	0.98	0.97	0.97	413
Mediterranean	0.97	0.99	0.98	308
Paleo	0.99	0.99	0.99	279
accuracy			0.98	1000
macro avg	0.98	0.98	0.98	1000
weighted avg	0.98	0.98	0.98	1000

CV Accuracy: 97.86% ± 0.50%

---

Model: SVM

Accuracy: 98.9 %

F1 Score: 0.989

## Classification Report:

	precision	recall	f1-score	support
Keto	1.00	0.98	0.99	413
Mediterranean	0.98	0.99	0.99	308
Paleo	0.98	1.00	0.99	279
accuracy			0.99	1000
macro avg	0.99	0.99	0.99	1000
weighted avg	0.99	0.99	0.99	1000

CV Accuracy: 99.01% ± 0.20%

---

Model: LightGBM

Accuracy: 99.4 %

F1 Score: 0.994

## Classification Report:

	precision	recall	f1-score	support
Keto	1.00	0.99	1.00	413
Mediterranean	0.98	1.00	0.99	308
Paleo	1.00	0.99	0.99	279
accuracy			0.99	1000
macro avg	0.99	0.99	0.99	1000
weighted avg	0.99	0.99	0.99	1000

CV Accuracy: 99.76% ± 0.25%

---

Model: CatBoost

Accuracy: 99.6 %

F1 Score: 0.996

## Classification Report:

	precision	recall	f1-score	support
Keto	1.00	0.99	1.00	413
Mediterranean	0.99	1.00	1.00	308
Paleo	1.00	1.00	1.00	279
accuracy			1.00	1000
macro avg	1.00	1.00	1.00	1000
weighted avg	1.00	1.00	1.00	1000

CV Accuracy: 99.62% ± 0.20%

---

In [653]:

```
# ----- Imports -----
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
```

```

from collections import Counter
import warnings
warnings.filterwarnings("ignore")

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import accuracy_score, f1_score, classification_report, confusion_matrix

from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from lightgbm import LGBMClassifier
from imblearn.over_sampling import SMOTE

# ----- Preprocessing -----

# Add Gaussian noise to reduce overfit
np.random.seed(42)
df['Carbohydrate_Intake'] += np.random.normal(0, 20, df.shape[0])
df['Fat_Intake'] += np.random.normal(0, 15, df.shape[0])

# Create target label
df['Smart_Diet'] = df.apply(
    lambda row: 'Low Carb' if row['Carbohydrate_Intake'] < 180 else
                'High Fat' if row['Fat_Intake'] > 90 else
                'Balanced', axis=1)
df['Recommended_Meal_Plan'] = df['Smart_Diet'].map({
    'Low Carb': 'Keto',
    'High Fat': 'Paleo',
    'Balanced': 'Mediterranean'
})
df.drop(columns=['Smart_Diet'], inplace=True)

# Drop unnecessary columns
df.drop(columns=[col for col in ['Patient_ID', 'BMI'] if col in df.columns], inplace=True)

# Handle missing values
imputer = SimpleImputer(strategy='constant', fill_value='Unknown')
for col in ['Chronic_Disease', 'Allergies', 'Food_Aversions']:
    if col in df.columns and df[col].dtype == 'object':
        df[col] = imputer.fit_transform(df[[col]]).ravel()

# Encode categorical features
cat_cols = df.select_dtypes(include='object').columns.difference(['Recommended_Meal_Plan'])
encoders = {}
for col in cat_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    encoders[col] = le

# Encode target
le_rmp = LabelEncoder()
df['Recommended_Meal_Plan'] = le_rmp.fit_transform(df['Recommended_Meal_Plan'])

# Split features and target
X = df.drop(columns=['Recommended_Meal_Plan'])
y = df['Recommended_Meal_Plan']

X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)

# ----- SMOTE Balancing -----

print("Original training class distribution:", Counter(y_train))

# Get max count among classes
max_count = max(Counter(y_train).values())

```

```

# Create sampling strategy to oversample minority classes
sampling_strategy = {cls: max_count for cls in np.unique(y_train)}

# Apply SMOTE
smote = SMOTE(sampling_strategy=sampling_strategy, random_state=42)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)

print("After SMOTE:", Counter(y_train_bal))

# Add noise to reduce model memorization
X_train_bal += np.random.normal(0, 0.01, X_train_bal.shape)

# ----- Scaling -----

scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_bal)
X_test_scaled = scaler.transform(X_test)

# ----- Feature Selection -----

rf = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)
rf.fit(X_train_scaled, y_train_bal)
selector = SelectFromModel(rf, threshold="1.0*mean", prefit=True)

X_train_sel = selector.transform(X_train_scaled)
X_test_sel = selector.transform(X_test_scaled)
selected_features = X.columns[selector.get_support()].tolist()
print("Selected Features:", selected_features)

# ----- Model Setup -----

models = {
    "SVM": SVC(
        kernel='rbf', C=1.0, gamma='scale', class_weight='balanced',
        probability=True, random_state=42
    ),
}

# ----- Evaluation -----

for name, model in models.items():
    print(f"\nModel: {name}")
    model.fit(X_train_sel, y_train_bal)
    y_pred = model.predict(X_test_sel)

    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred, average='weighted')

    print("Accuracy:", round(acc * 100, 2), "%")
    print("F1 Score:", round(f1, 4))
    print("Classification Report:\n", classification_report(y_test, y_pred, target_names=
le_rmp.classes_))

    cv = cross_val_score(model, X_train_sel, y_train_bal, cv=5, scoring='accuracy')
    print(f"CV Accuracy: {cv.mean() * 100:.2f}% ± {cv.std() * 100:.2f}%")
    print("-" * 80)

```

Original training class distribution: Counter({1: 1481, 2: 1318, 0: 1201})  
After SMOTE: Counter({0: 1481, 1: 1481, 2: 1481})  
Selected Features: ['Carbohydrate\_Intake', 'Fat\_Intake', 'Recommended\_Carbs']

Model: SVM

Accuracy: 98.4 %

F1 Score: 0.984

Classification Report:

	precision	recall	f1-score	support
Keto	0.97	0.99	0.98	301
Mediterranean	1.00	0.97	0.98	370
Paleo	0.98	0.99	0.99	329
accuracy			0.98	1000

macro avg	0.98	0.98	0.98	1000
weighted avg	0.98	0.98	0.98	1000

CV Accuracy: 98.38% ± 0.52%

---

In [654]:

```
# ----- Imports -----
import pandas as pd
import numpy as np
from collections import Counter
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import accuracy_score, f1_score, classification_report
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from lightgbm import LGBMClassifier
from imblearn.over_sampling import SMOTE
import warnings
warnings.filterwarnings("ignore")

# ----- Data Preprocessing -----

# Add slight noise to reduce overfit before modeling
np.random.seed(42)
df['Carbohydrate_Intake'] += np.random.normal(0, 10, df.shape[0])
df['Fat_Intake'] += np.random.normal(0, 8, df.shape[0])

# Generate target
df['Smart_Diet'] = df.apply(
    lambda row: 'Low Carb' if row['Carbohydrate_Intake'] < 180 else
                'High Fat' if row['Fat_Intake'] > 90 else
                'Balanced', axis=1)
df['Recommended_Meal_Plan'] = df['Smart_Diet'].map({
    'Low Carb': 'Keto',
    'High Fat': 'Paleo',
    'Balanced': 'Mediterranean'
})
df.drop(columns=['Smart_Diet'], inplace=True)

# Drop irrelevant columns
df.drop(columns=[col for col in ['Patient_ID', 'BMI'] if col in df.columns], inplace=True)

# Fill missing values
imputer = SimpleImputer(strategy='constant', fill_value='Unknown')
for col in ['Chronic_Disease', 'Allergies', 'Food_Aversions']:
    if col in df.columns and df[col].dtype == 'object':
        df[col] = imputer.fit_transform(df[[col]]).ravel()

# Encode categorical variables
cat_cols = df.select_dtypes(include='object').columns.difference(['Recommended_Meal_Plan'])
encoders = {}
for col in cat_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    encoders[col] = le

# Encode target variable
le_rmp = LabelEncoder()
df['Recommended_Meal_Plan'] = le_rmp.fit_transform(df['Recommended_Meal_Plan'])

# ----- Train-Test Split -----
X = df.drop(columns=['Recommended_Meal_Plan'])
y = df['Recommended_Meal_Plan']
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)
```

```

# ----- SMOTE Balancing -----
print("Original class distribution:", Counter(y_train))
max_count = max(Counter(y_train).values())
smote = SMOTE(sampling_strategy={cls: max_count for cls in np.unique(y_train)}, random_s
tate=42)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)
print("After SMOTE:", Counter(y_train_bal))

# ----- Scaling -----
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_bal)
X_test_scaled = scaler.transform(X_test)

# ----- Feature Selection -----
rf = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)
rf.fit(X_train_scaled, y_train_bal)
selector = SelectFromModel(rf, threshold="0.9*mean", prefit=True) # slightly relaxed thr
eshold
X_train_sel = selector.transform(X_train_scaled)
X_test_sel = selector.transform(X_test_scaled)
selected_features = X.columns[selector.get_support()].tolist()
print("Selected Features:", selected_features)

# ----- Models -----
models = {
    "SVM": SVC(
        kernel='rbf', C=1.0, gamma='scale',
        class_weight='balanced', probability=True, random_state=42
    ),
    "LightGBM": LGBMClassifier(
        num_leaves=20,
        max_depth=4,
        learning_rate=0.05,
        n_estimators=100,
        subsample=0.7,
        colsample_bytree=0.7,
        reg_lambda=5.0,
        reg_alpha=3.0,
        random_state=42,
        verbose=-1 # <--- Suppresses all LightGBM internal training logs
    )
}

# ----- Evaluation -----
for name, model in models.items():
    print(f"\n===== {name} Evaluation =====")
    model.fit(X_train_sel, y_train_bal)
    y_pred = model.predict(X_test_sel)

    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred, average='weighted')
    report = classification_report(y_test, y_pred, target_names=le_rmp.classes_)

    print(f"Accuracy: {acc * 100:.2f}%")
    print(f"F1 Score: {f1:.4f}")
    print("Classification Report:")
    print(report)

    cv = cross_val_score(model, X_train_sel, y_train_bal, cv=5, scoring='accuracy')
    print(f"Cross-Val Accuracy: {cv.mean() * 100:.2f}% ± {cv.std() * 100:.2f}%")
    print("=" * 60)

```

Original class distribution: Counter({1: 1476, 2: 1298, 0: 1226})  
 After SMOTE: Counter({0: 1476, 1: 1476, 2: 1476})  
 Selected Features: ['Carbohydrate\_Intake', 'Fat\_Intake', 'Recommended\_Carbs']

==== SVM Evaluation =====  
 Accuracy: 98.80%  
 F1 Score: 0.9880  
 Classification Report:

	precision	recall	f1-score	support
Keto	0.98	0.99	0.99	306
Mediterranean	0.99	0.99	0.99	369
Paleo	1.00	0.98	0.99	325
accuracy			0.99	1000
macro avg	0.99	0.99	0.99	1000
weighted avg	0.99	0.99	0.99	1000

Cross-Val Accuracy: 98.35%  $\pm$  0.32%

==== LightGBM Evaluation =====

Accuracy: 99.70%

F1 Score: 0.9970

Classification Report:

	precision	recall	f1-score	support
Keto	1.00	1.00	1.00	306
Mediterranean	0.99	1.00	1.00	369
Paleo	1.00	0.99	1.00	325
accuracy			1.00	1000
macro avg	1.00	1.00	1.00	1000
weighted avg	1.00	1.00	1.00	1000

Cross-Val Accuracy: 99.59%  $\pm$  0.23%

In [695]:

```
# ----- Imports -----
import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
from collections import Counter
import warnings
warnings.filterwarnings("ignore")

from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.impute import SimpleImputer
from sklearn.feature_selection import SelectFromModel
from sklearn.metrics import accuracy_score, f1_score, classification_report, confusion_matrix

from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from lightgbm import LGBMClassifier
from imblearn.over_sampling import SMOTE

# ----- Load & Add Noise -----
np.random.seed(42)
df['Carbohydrate_Intake'] += np.random.normal(0, 20, df.shape[0])
df['Fat_Intake'] += np.random.normal(0, 15, df.shape[0])

# ----- Create Targets -----
df['Smart_Diet'] = df.apply(
    lambda row: 'Low Carb' if row['Carbohydrate_Intake'] < 180 else
                'High Fat' if row['Fat_Intake'] > 90 else
                'Balanced', axis=1)

df['Recommended_Meal_Plan'] = df['Smart_Diet'].map({
    'Low Carb': 'Keto',
    'High Fat': 'Paleo',
    'Balanced': 'Mediterranean'
})
df.drop(columns=['Smart_Diet'], inplace=True)

# ----- Handle Missing Values -----
```

```

imputer = SimpleImputer(strategy='constant', fill_value='Unknown')
for col in ['Chronic_Disease', 'Allergies', 'Food_Aversions']:
    if col in df.columns and df[col].dtype == 'object':
        df[col] = imputer.fit_transform(df[[col]]).ravel()

# ----- Feature Engineering -----
def compute_health_score(row):
    score = 0
    if row['Blood_Pressure_Systolic'] >= 140 or row['Blood_Pressure_Diastolic'] >= 90:
        score += 2
    elif row['Blood_Pressure_Systolic'] <= 90 or row['Blood_Pressure_Diastolic'] <= 60:
        score += 1
    if row['Cholesterol_Level'] == 'High':
        score += 2
    if row['Blood_Sugar_Level'] == 'High':
        score += 2
    return score

def compute_risk_profile(row):
    score = 0
    if row['Chronic_Disease'] != 'None':
        score += 2
    if row['Genetic_Risk_Factor'] != 'None':
        score += 2
    if row['Smoking_Habit'] != 'Non-Smoker':
        score += 1
    if row['Alcohol_Consumption'] != 'None':
        score += 1
    if 'BMI' in df.columns and row['BMI'] > 30:
        score += 2
    return score

df['Health_Score'] = df.apply(compute_health_score, axis=1)
df['Risk_Score'] = df.apply(compute_risk_profile, axis=1)

# ----- Drop Unnecessary Columns -----
df.drop(columns=[col for col in ['Patient_ID', 'BMI'] if col in df.columns], inplace=True)

# ----- Encode Categorical Features -----
cat_cols = df.select_dtypes(include='object').columns.difference(['Recommended_Meal_Plan'])
encoders = {}
for col in cat_cols:
    le = LabelEncoder()
    df[col] = le.fit_transform(df[col])
    encoders[col] = le

# Encode target
le_rmp = LabelEncoder()
df['Recommended_Meal_Plan'] = le_rmp.fit_transform(df['Recommended_Meal_Plan'])

# ----- Train-Test Split -----
X = df.drop(columns=['Recommended_Meal_Plan'])
y = df['Recommended_Meal_Plan']
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)

# ----- SMOTE Balancing -----
print("Original training class distribution:", Counter(y_train))

max_count = max(Counter(y_train).values())
sampling_strategy = {cls: max_count for cls in np.unique(y_train)}

smote = SMOTE(sampling_strategy=sampling_strategy, random_state=42)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)

print("After SMOTE:", Counter(y_train_bal))

# Add small noise to prevent memorization
X_train_bal += np.random.normal(0, 0.01, X_train_bal.shape)

```

```

# ----- Scaling -----
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_bal)
X_test_scaled = scaler.transform(X_test)

# ----- Feature Selection -----
rf = RandomForestClassifier(n_estimators=100, max_depth=10, random_state=42)
rf.fit(X_train_scaled, y_train_bal)
selector = SelectFromModel(rf, threshold="1.0*mean", prefit=True)

X_train_sel = selector.transform(X_train_scaled)
X_test_sel = selector.transform(X_test_scaled)

selected_features = X.columns[selector.get_support()].tolist()
print("Selected Features:", selected_features)

# ----- Model Setup -----
models = {
    "SVM": SVC(
        kernel='rbf', C=1.0, gamma='scale', class_weight='balanced',
        probability=True, random_state=42
    ),
}

# ----- Evaluation -----
for name, model in models.items():
    print(f"\nModel: {name}")
    model.fit(X_train_sel, y_train_bal)
    y_pred = model.predict(X_test_sel)

    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred, average='weighted')

    print("Accuracy:", round(acc * 100, 2), "%")
    print("F1 Score:", round(f1, 4))
    print("Classification Report:\n", classification_report(y_test, y_pred, target_names=
le_rmp.classes_))

    cv = cross_val_score(model, X_train_sel, y_train_bal, cv=5, scoring='accuracy')
    print(f"CV Accuracy: {cv.mean() * 100:.2f}% ± {cv.std() * 100:.2f}%")
    print("-" * 80)

```

Original training class distribution: Counter({0: 1677, 1: 1223, 2: 1100})

After SMOTE: Counter({1: 1677, 0: 1677, 2: 1677})

Selected Features: ['Carbohydrate\_Intake', 'Fat\_Intake']

Model: SVM

Accuracy: 98.8 %

F1 Score: 0.988

Classification Report:

	precision	recall	f1-score	support
Keto	1.00	0.98	0.99	419
Mediterranean	0.98	1.00	0.99	306
Paleo	0.99	0.99	0.99	275
accuracy			0.99	1000
macro avg	0.99	0.99	0.99	1000
weighted avg	0.99	0.99	0.99	1000

CV Accuracy: 98.77% ± 0.47%

In [656]:

```

# ----- Imports -----
import pandas as pd
import numpy as np
import warnings
from collections import Counter
from sklearn.model_selection import train_test_split, cross_val_score
from sklearn.preprocessing import LabelEncoder, StandardScaler

```



```

from sklearn.impute import SimpleImputer
from sklearn.metrics import accuracy_score, f1_score, classification_report
from sklearn.ensemble import RandomForestClassifier
from sklearn.svm import SVC
from lightgbm import LGBMClassifier
from imblearn.over_sampling import SMOTE

warnings.filterwarnings("ignore")

# ----- Step 1: Add Noise -----
np.random.seed(42)
df['Carbohydrate_Intake'] += np.random.normal(0, 20, df.shape[0])
df['Fat_Intake'] += np.random.normal(0, 15, df.shape[0])

# ----- Step 2: Generate Targets -----
df['Smart_Diet'] = df.apply(
    lambda row: 'Low Carb' if row['Carbohydrate_Intake'] < 180 else
                'High Fat' if row['Fat_Intake'] > 90 else
                'Balanced', axis=1)

df['Recommended_Meal_Plan'] = df['Smart_Diet'].map({
    'Low Carb': 'Keto',
    'High Fat': 'Paleo',
    'Balanced': 'Mediterranean'
})
df.drop(columns=['Smart_Diet'], inplace=True)

# ----- Step 3: Missing Values -----
imputer = SimpleImputer(strategy='constant', fill_value='Unknown')
for col in ['Chronic_Disease', 'Allergies', 'Food_Aversions']:
    if col in df.columns and df[col].dtype == 'object':
        df[col] = imputer.fit_transform(df[[col]]).ravel()

# ----- Step 4: Feature Engineering -----
def compute_health_score(row):
    score = 0
    if row['Blood_Pressure_Systolic'] >= 140 or row['Blood_Pressure_Diastolic'] >= 90:
        score += 2
    elif row['Blood_Pressure_Systolic'] <= 90 or row['Blood_Pressure_Diastolic'] <= 60:
        score += 1
    if row['Cholesterol_Level'] == 'High':
        score += 2
    if row['Blood_Sugar_Level'] == 'High':
        score += 2
    return score

def compute_risk_profile(row):
    score = 0
    if row['Chronic_Disease'] != 'None':
        score += 2
    if row['Genetic_Risk_Factor'] != 'None':
        score += 2
    if row['Smoking_Habit'] != 'Non-Smoker':
        score += 1
    if row['Alcohol_Consumption'] != 'None':
        score += 1
    if 'BMI' in df.columns and row['BMI'] > 30:
        score += 2
    return score

df['Health_Score'] = df.apply(compute_health_score, axis=1)
df['Risk_Score'] = df.apply(compute_risk_profile, axis=1)

# ----- Step 5: Drop Unnecessary -----
df.drop(columns=[col for col in ['Patient_ID', 'BMI'] if col in df.columns], inplace=True)

# ----- Step 6: Encode -----
cat_cols = df.select_dtypes(include='object').columns.difference(['Recommended_Meal_Plan'])
for col in cat_cols:
    le = LabelEncoder()

```

```

df[col] = le.fit_transform(df[col])

le_rmp = LabelEncoder()
df['Recommended_Meal_Plan'] = le_rmp.fit_transform(df['Recommended_Meal_Plan'])

# ----- Step 7: Split -----
X = df.drop(columns=['Recommended_Meal_Plan'])
y = df['Recommended_Meal_Plan']
X_train, X_test, y_train, y_test = train_test_split(X, y, stratify=y, test_size=0.2, random_state=42)

# ----- Step 8: Standard Scaling -----
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train)
X_test_scaled = scaler.transform(X_test)

# ----- Step 9: Train Models (Before SMOTE) -----
models = {
    "Random Forest": RandomForestClassifier(n_estimators=100, random_state=42),
    "SVM": SVC(kernel='rbf', C=1.0, gamma='scale', probability=True, class_weight='balanced', random_state=42),
    "LightGBM": LGBMClassifier(random_state=42)
}

print("\n Results BEFORE SMOTE")
for name, model in models.items():
    model.fit(X_train_scaled, y_train)
    y_pred = model.predict(X_test_scaled)

    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred, average='weighted')

    print(f"\nModel: {name}")
    print("Accuracy:", round(acc * 100, 2), "%")
    print("F1 Score:", round(f1, 4))
    print("-" * 50)

# ----- Step 10: Apply SMOTE -----
print("\nOriginal training distribution:", Counter(y_train))

smote = SMOTE(random_state=42)
X_train_bal, y_train_bal = smote.fit_resample(X_train, y_train)
print("After SMOTE distribution:", Counter(y_train_bal))

X_train_bal_scaled = scaler.fit_transform(X_train_bal)

# ----- Step 11: Train Models (After SMOTE) -----
print("\n Results AFTER SMOTE")
for name, model in models.items():
    model.fit(X_train_bal_scaled, y_train_bal)
    y_pred = model.predict(X_test_scaled)

    acc = accuracy_score(y_test, y_pred)
    f1 = f1_score(y_test, y_pred, average='weighted')

    print(f"\nModel: {name}")
    print("Accuracy:", round(acc * 100, 2), "%")
    print("F1 Score:", round(f1, 4))
    print("-" * 50)

```

Results BEFORE SMOTE

Model: Random Forest  
Accuracy: 100.0 %  
F1 Score: 1.0

-----

Model: SVM  
Accuracy: 92.1 %  
F1 Score: 0.9209

-----

```
Model: LightGBM
Accuracy: 99.7 %
F1 Score: 0.997
-----
```

```
Original training distribution: Counter({1: 1396, 0: 1350, 2: 1254})
After SMOTE distribution: Counter({1: 1396, 0: 1396, 2: 1396})
```

Results AFTER SMOTE

```
Model: Random Forest
Accuracy: 99.0 %
F1 Score: 0.99
-----
```

```
Model: SVM
Accuracy: 92.0 %
F1 Score: 0.92
-----
```

```
Model: LightGBM
Accuracy: 98.9 %
F1 Score: 0.989
-----
```

In [656]:

